

Pandas


CS3300 Data Science

RJ Nowling

Readings

- Chapters 2 and 3 of the *Python Data Science Handbook*

Data Science Process



What is Pandas?

- Library for manipulating tables of data
- Primarily used for cleaning and restructuring data in preparation for plotting or modeling
- 3 primary data structures
 - Series – 1D, columns of data
 - DataFrames – 2D, tables of data
 - Panel – 3D, cube of data (rarely used, deprecated and going to be removed)
- Columnar
 - Most operations are designed to operate on columns of data, not individual elements or rows

Pandas vs Numpy

Numpy

- Any dimension
- Indexing by position (e.g., row or column)
- Usually a single type (e.g., int, float)

Pandas

- Limited to 1 (Series) or 2 (DataFrame) dimensions
- Indexing primarily by column names
- Each column has a its own type

Pandas and SQL

- A DataFrame is similar to a table in a SQL database
- SQL also operates on columns
- Many Pandas operations have analogs in SQL:
 - Head – `select * from <table> limit 10;`
 - Selecting columns – `select column1, column2, ... from <table>;`
 - Filtering rows – `where column1 > 5;`
 - Grouping rows – `select max(...), column2 ... group by column2;`
 - Joining tables – `select ... from table1 join table2 on table1.column1 = table2.column1;`
- Pandas DataFrames have an index – this is normally the implicit numerical index

Caveats

- Pandas offers multiple ways to do things. Some ways are newer and have learned from the mistakes of the old ways. This can be confusing and frustrating
- The pandas documentation is complex and not well organized
- It can be difficult to predict when a copy is made versus a view is created – this makes optimization challenging

Creating DataFrames

- Read from a CSV file:

```
df = pd.read_csv(filename)
```

- From existing lists, Numpy, arrays, or series:

```
df = pd.DataFrame( { "column1" : [0.0, 1.0, 2.0],  
                    "column2" : np_array,  
                    "column3" : series } )
```


Investigating DataFrames

- `df.head()`
- `df.dtypes`
- `df.shape`

Indexing / Selecting Columns

- Pandas has multiple ways to index. The slice operator works on columns:

`df["column name"]` – get a single column

`df[["column1", "column2", "column3"]]` – get multiple columns

Indexing

- You can index by position (numerical index). This follows the Numpy pattern of row, then column:

`df.iloc[5]` – get a single row

`df.iloc[:, :3]` – get first 3 columns and all rows of the df

`df.iloc[1:100]` – get the second to 100th rows

Creating a New Column

- The simplest way to create a new column:

```
df["new column name"] = <list, 1D numpy array,  
Series>
```

- The assign method is useful since it returns a new DataFrame and can be used with method chaining:

```
new_df = df.assign(<new_column_name> = <list,  
1D numpy array, Series>)
```

Modifying a Column

- Convert data types – may need to specify function for parsing / conversion
- Cleaning data
- Extracting fields from complex types
 - e.g., hour, month, etc. from date times

Modifying a Column

1. Get the Series for the column of interest

```
column = df["column name"]
```

2. Use the map() method to apply a function to each element in the Series and return a new Series

```
converted = column.map(lambda s: s + 1)
```

3. Then update the df, either by adding a new column or overwriting the original column

```
df["column name"] = converted
```

Dropping a Column

- I prefer to use the `drop()` method because it returns a DataFrame object so it works with chaining:



```
new_df = df.drop(columns=["column name"])
```

- You might see this, too:

```
del df["column_name"]
```

I like this less because it updates the DataFrame in place.


Inner Joins




key	val_x	val_y
1	x1	y1
2	x2	y2

Left Outer Join

LEFT JOIN





Left



key	val_x	val_y
1	x1	y1
2	x2	y2
3	x3	NA


Right Outer Join

RIGHT JOIN



key	val_x	val_y
1	x1	y1
2	x2	y2
4	NA	y3

Full Outer Join



Pandas Join

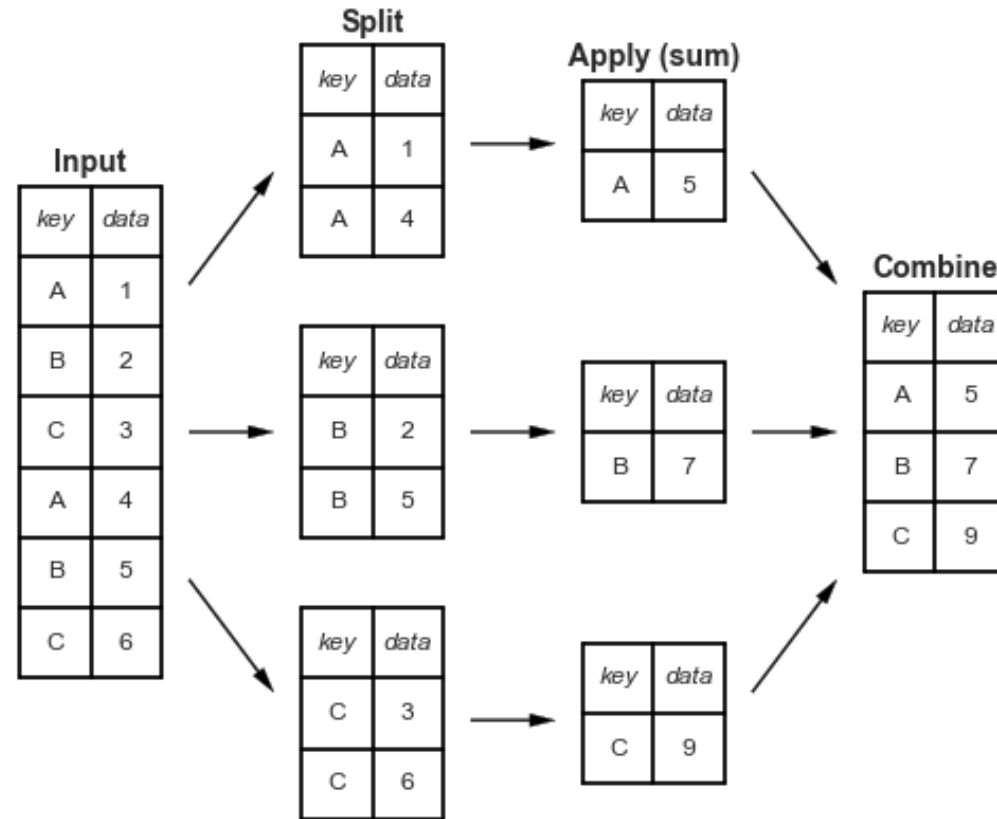
```
joined_df = df1.merge(df2,  
                      on = "column",  
                      how = "inner")
```

```
joined_df = df1.merge(df2,  
                      left_on = "column1",  
                      right_on = "column2",  
                      how = "outer")
```

Group By and Aggregations

- Group by groups columns by a set of keys and aggregates the values in the remaining columns
- Two types of columns participate in group by:
 - Columns that are the keys – the keys have to have the exact same values for rows to be matched up
 - Aggregated columns – we apply count, min, max, sum, or a similar function to reduce multiple values to a single value

Group By and Aggregations



Group By and Aggregations

My preferred recipe:

1. Select the columns you want to use:

```
df_to_group = df[["column1", "column2", "column3"]]
```

2. Perform the group by:

```
grouped_df = df_to_group.groupby(by = ["column1",  
                                       "column2"],  
                                as_index = False)  
                                .min()
```